



Version 4

04/18/21

Provided by Tech Skills Transformations LLC



Prerequisites

- Modern system with at least 8G of memory and 20G free storage
- VirtualBox installed and running

<http://www.virtualbox.org>

- Virtual machine (.ova file) installed in VirtualBox

<https://www.dropbox.com/s/konq5fctnalu9lz/k8s-op.ova?dl=0>

OR

<https://bclconf.s3-us-west-2.amazonaws.com/k8s-op.ova>

- Workshop docs are in <https://github.com/skilldocs/k8s-op>

- Setup doc is at

<https://github.com/skilldocs/k8s-op/blob/main/k8s-op-setup.pdf>

- Labs doc for workshop

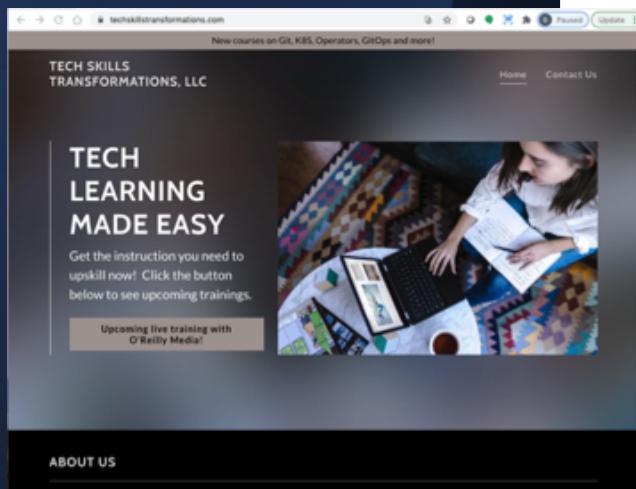
<https://github.com/skilldocs/k8s-op/blob/main/k8s-op-labs.pdf>



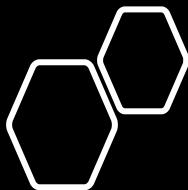
Building a Kubernetes Operator: Extending Kubernetes to Fit your Applications

Brent Laster

About me



- Founder, Tech Skills Transformations LLC
- R&D DevOps Director
- Global trainer – training (Git, Jenkins, Gradle, CI/CD, pipelines, Kubernetes, Helm, ArgoCD, operators)
- Author -
 - OpenSource.com
 - Professional Git book
 - Jenkins 2 – Up and Running book
 - Continuous Integration vs. Continuous Delivery vs. Continuous Deployment mini-book on Safari
- <https://www.linkedin.com/in/brentlaster>
- @BrentCLaster
- GitHub: brentlaster



Professional Git Book

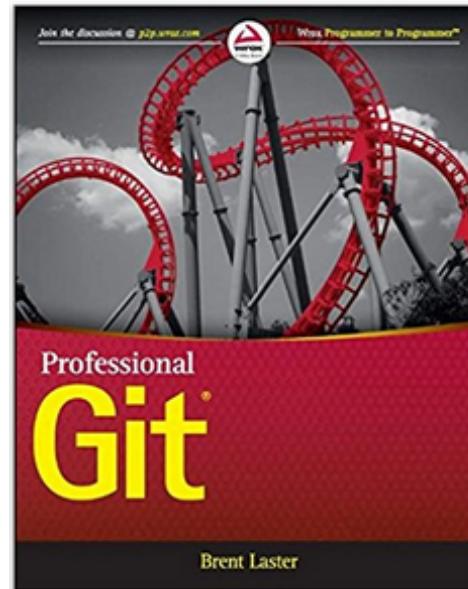
- Extensive Git reference, explanations,
- and examples
- First part for non-technical
- Beginner and advanced reference
- Hands-on labs

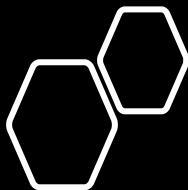
Professional Git 1st Edition

by [Brent Laster](#) (Author)

7 customer reviews

[Look inside](#)





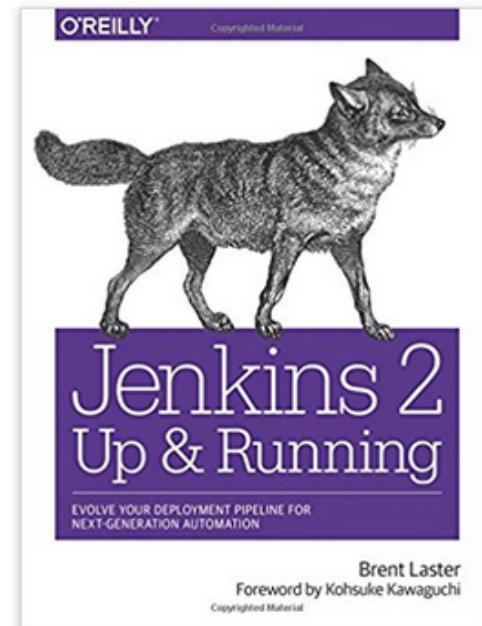
Jenkins 2 Book

- Jenkins 2 – Up and Running
- “It’s an ideal book for those who are new to CI/CD, as well as those who have been using Jenkins for many years. This book will help you discover and rediscover Jenkins.” *By Kohsuke Kawaguchi, Creator of Jenkins*

★★★★★ 5 customer reviews

#1 New Release in Java Programming

[Look inside](#) ↴





O'Reilly Training

April 26 & 27, 2021

Git Fundamentals

Join Brent Laster to gain the knowledge and skills you need to leverage Git to greatly simplify and speed up managing all of the changes in your source code.

Once you ...

LIVE ONLINE TRAINING

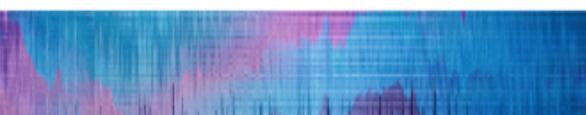
Next Level Git - Master your content

Use powerful tools in Git to simplify merges, rewrite history, and perform automatic updates

Topic: Software Development



BRENT LASTER



LIVE ONLINE TRAINING

Building a deployment pipeline with Jenkins 2

Manage continuous integration and continuous delivery to release software

Topic: System Administration



BRENT LASTER



LIVE ONLINE TRAINING

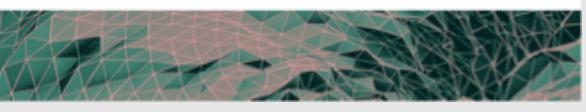
Containers A-Z

An overview of containers, Docker, Kubernetes, Istio, Helm, Kubernetes Operators, and GitOps

Topic: System Administration



BRENT LASTER



June 3, 10, 17 & 24, 2021

Git in 4 Weeks

Join author, trainer, and DevOps director Brent Laster to learn how Git works—and discover all the things you can do with it. Over four sessions, Brent walks you through everything you ...

LIVE ONLINE TRAINING

Next Level Git - Master your workflow

Use Git to find problems, simplify working with multiple branches and repositories, and customize behavior with hooks

Topic: Software Development



BRENT LASTER



LIVE ONLINE TRAINING

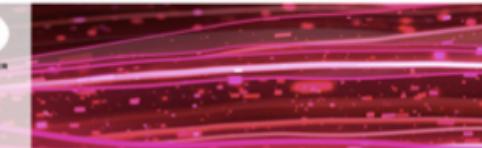
Git Troubleshooting

How to solve practically any problem that comes your way

Topic: Software Development



BRENT LASTER



LIVE ONLINE TRAINING

Getting started with continuous delivery (CD)

Move beyond CI to build, manage, and deploy a working pipeline

Topic: System Administration



BRENT LASTER



LIVE ONLINE TRAINING

Helm Fundamentals

Deploying, upgrading, and rolling back applications in Kubernetes

Topic: System Administration



BRENT LASTER



June 28, 2021

Troubleshooting Kubernetes

In this 3-hour course, global trainer, author, and DevOps director Brent Laster will show you how to respond to the most common problem situations you may encounter with Kubernetes. You'll learn ...



May 24, 2021

Continuous Delivery in Kubernetes with ArgoCD

Join expert Brent Laster to explore GitOps and learn how to use Argo CD to implement GitOps in your Kubernetes deployments. APAC time friendly - You're a Kubernetes admin who wants ...



May 17, 2021

Building a Kubernetes Operator

Join expert Brent Laster to learn how the Operator pattern helps address these kinds of situations by allowing you to create custom controllers that extend the functionality of the Kubernetes API ...

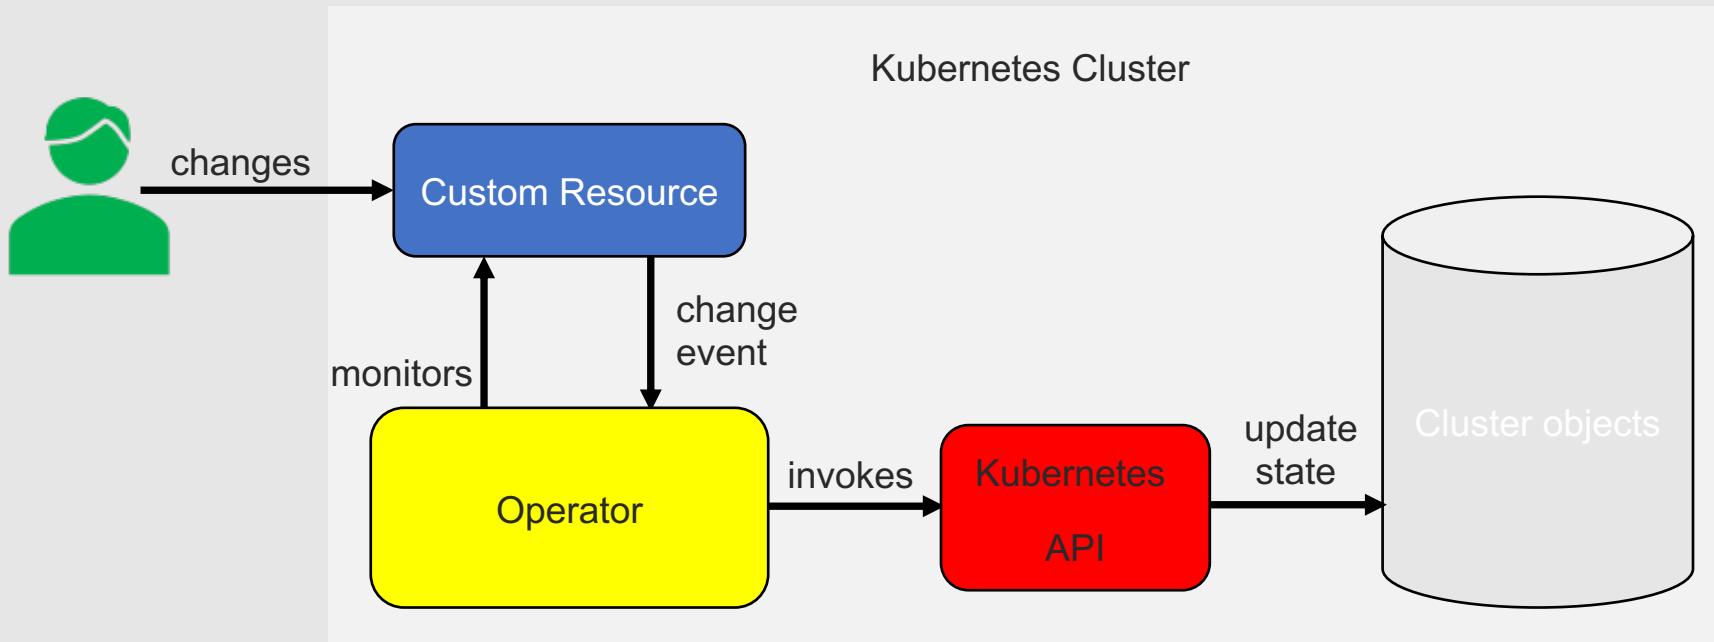
Agenda

- Introduction to Operators
- The Operator Framework and SDK
- Custom Resources
- All about Controllers
- Permissions and RBAC
- Initial Testing
- Packaging and Deploying the Operator in a Cluster



What is an Operator?

- A pattern for software extension to Kubernetes
- A way to package, deploy, and manage an application as K8s-native
 - K8s-native application - application deployed on K8s and managed via K8s-api and K8s tooling
- Application that watches a custom resource (new Kind)
- Executes some defined operations when custom resource changes





From a human perspective...

- A design pattern
 - Original blog by CoreOS from 2016 -
<https://coreos.com/blog/introducing-operators.html>
- Encapsulate operational knowledge in code
 - Integrates with Kubernetes and its API
 - “Kubernetes-native”
- Similar to an SRE role
- Automate/implement Day-1 activities
 - Installation, configuration, etc.
- Automate/implement Day-2 activities
 - Update, backup, failover, restore, re-configuration, etc.
- Manages applications that need to persist state (stateful)



Stateful vs Stateless

- Native k8s objects (deployments, services, etc.) are designed to work well with the standard k8s controller model
- Considered “stateless” – don’t need to persist/rely on state
- That idea has issues for more complex, **stateful** applications
 - Example: database running on multiple nodes
 - If nodes go down, need to reload data to get back to previous good state
- Usual cloud friendly operations like scaling, upgrading, disaster recovery for stateful app don’t fit



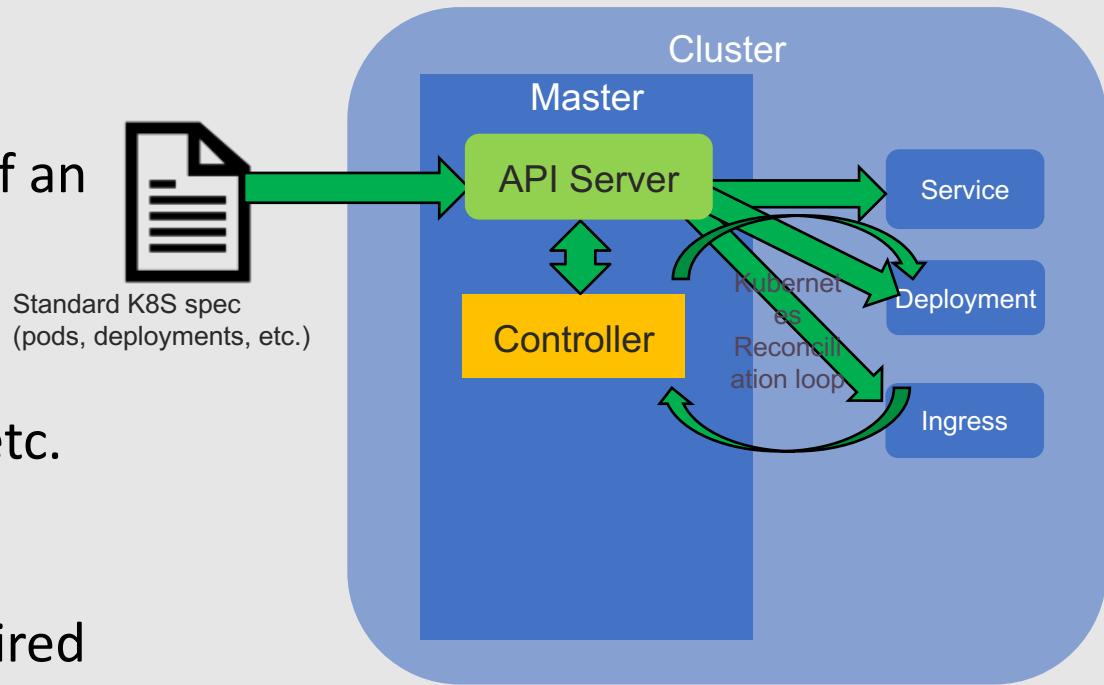
How are Operators different from other tools?

- Purpose-built to run a Kubernetes application, with operational knowledge baked in
- “Smarter” and more tailored than generic tools
- Cloud-like capabilities encoded into Operator code can provide an advanced user experience
- Can automate features like updates, backups and scaling
- All of this is accomplished using standard Kubernetes tools, CLI and API
- An operator is essentially like a custom controller



What is a controller?

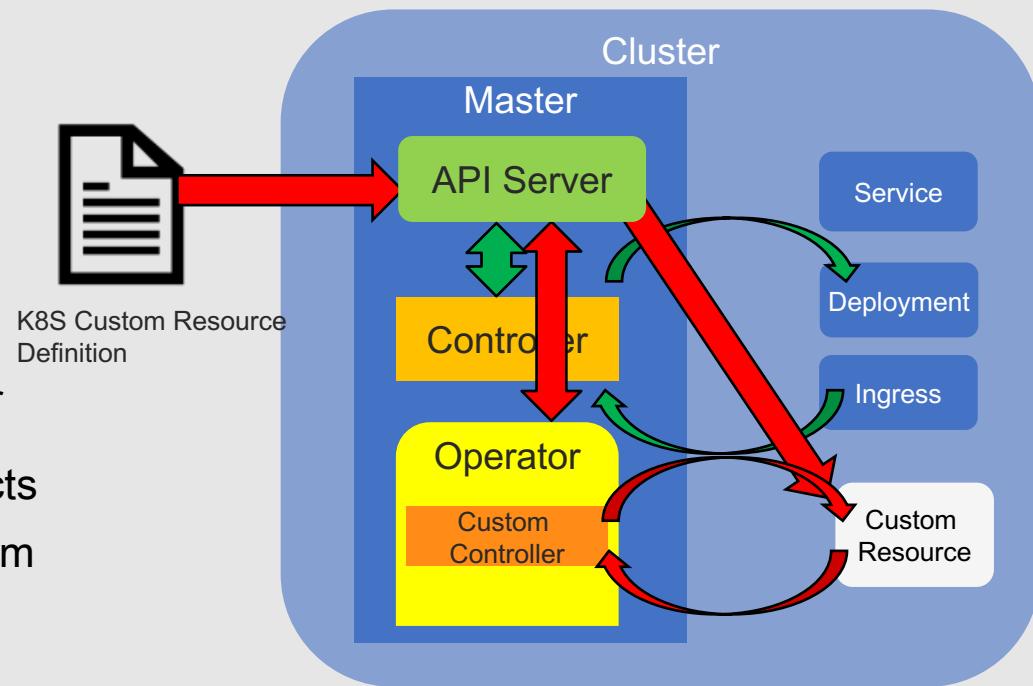
- Core part of Kubernetes
- Software loop that runs continuously on master node (in “control plane”)
 - Compares desired state and live/current state of an “object”
 - Objects are Kubernetes resources like Pods, ConfigMaps, Services, etc.
 - If differences in desired and live, controller reconciles to get to desired





How Do Operators Work?

- Define a custom controller to watch your app and do custom processing based on its state
- Application that is to be watched by operator is defined as new object type in k8s - custom resource (CR)
 - Has its own yaml spec and object type CustomResource
 - Format understood by the API server
- Operators run and oversee a custom controller loop
- Overall behavior of an operator's controller reconciling state for CR is similar to way native k8s handles reconciling native objects
- Operators give you the chance to do custom processing when reconciling





What can an Operator do?

- Simple to complex operations
- Operations on custom types
- Help manage stateful types
- Backup/restore
- Deploying clusters for a particular type
- Upgrading versions
- Additional logging
- Replication
- Creating additional resources
- Note: Could write operator to manage software that is not your own - if it doesn't already exist



When do you use an operator?

- Operators vs controllers
 - All operators are controllers but not all controllers are operators
 - To be considered as an operator, there must be application-specific domain knowledge built in
 - Able to perform automated tasks on behalf of user
- Operators are appropriate when
 - You need to encapsulate approach(es) for business logic around stateful applications (using Kubernetes API). Incorporates automation in Kubernetes for custom object.
 - You need to have an automated mechanism to watch your application for changes and respond with certain actions when certain events happen.
- Other alternative implementations may be appropriate/better



Alternatives to Operators*

- Outside of cluster
 - Shell scripts
 - Ansible automation
- Inside of cluster
 - Jobs
 - Custom controllers

* Often can be advisable to implement logic first in a simpler way to vet need for an operator



Writing an Operator

- Technically, could write it in any language/framework that can work with K8s API
- General approach
 - Code the logic for the operator – what do you want to do and how
 - Validate that an operator implementation makes sense!
 - » Are you compensating for architecture that needs to be changed?
 - Containerize the operator and make the image available in a registry
 - Create the Custom Resource Definition (CRD) – defines the custom resource
 - Create the RBAC definitions that will allow the operator to connect and work with resources in Kubernetes (via the API)
 - Package and deploy the operator into a cluster
 - Create the Custom Resource for it to monitor and respond to (based on the definition/structure provided by the CRD)
- Don't do everything from scratch - use available frameworks/tools



Operator Framework

WHAT'S IN THE FRAMEWORK?

Brief sentence about the contents of the framework to introduce the below websites. For example, the Operator Framework is a holistic framework for your Operator needs, etc.

BUILD

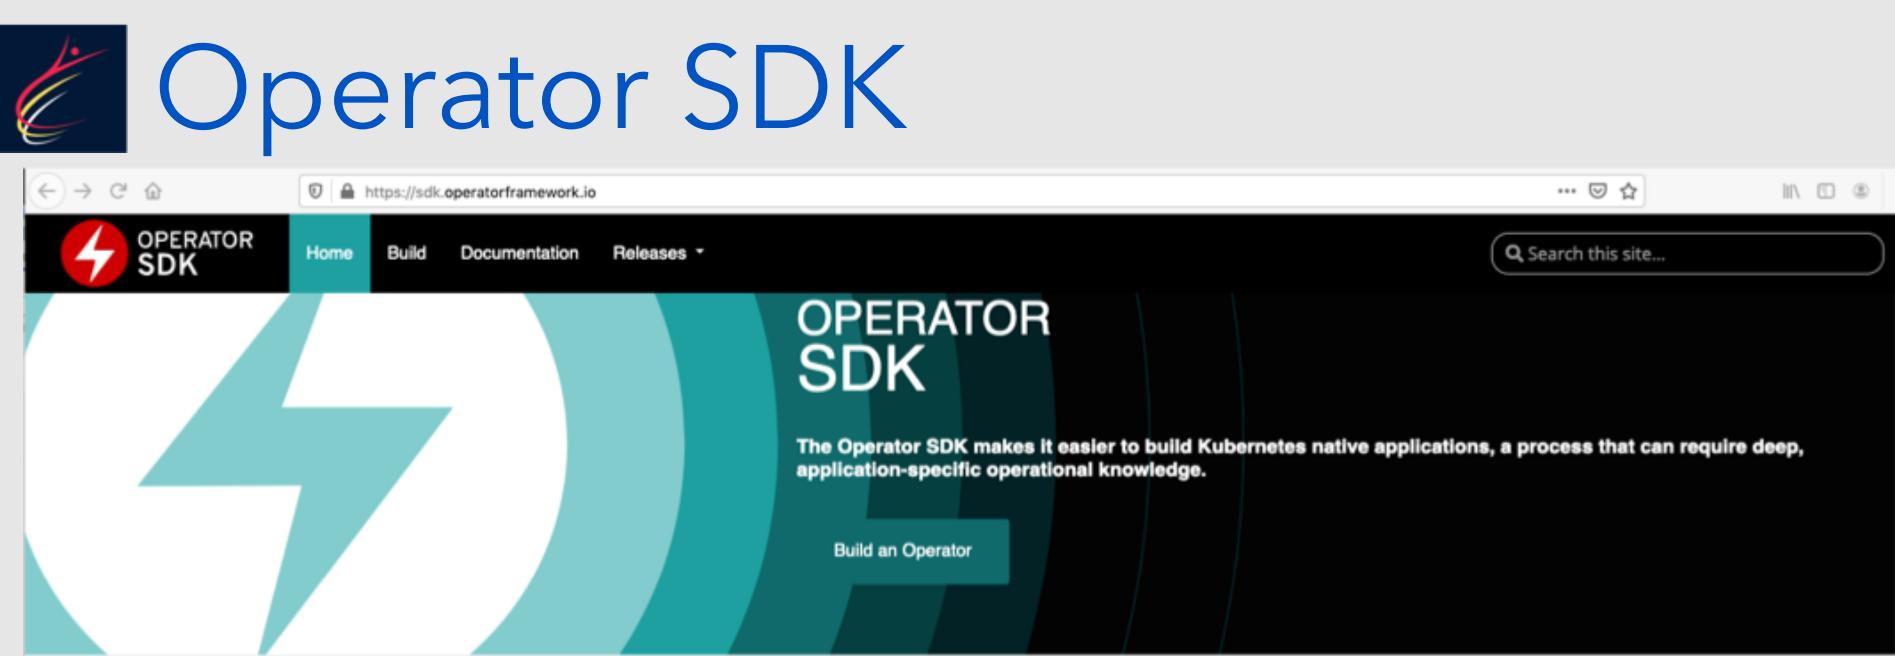
Software development kit for building Kubernetes applications.
Provides high level APIs, useful abstractions, and project scaffolding.
[Build an Operator →](#)

MANAGE

Management framework for extending Kubernetes with Operators.
Helps install, update, and manage the lifecycle of Operators running across clusters.
[Install OLM →](#)

DISCOVER

Home for the Kubernetes community to share Operators. Provides a catalog of existing Operators and guidance for contributing new Operators.
[Contribute an Operator →](#)



WHAT IS OPERATOR SDK?

This project is a component of the [Operator Framework](#), an open source toolkit to manage Kubernetes native applications, called Operators, in an effective, automated, and scalable way.

WHAT CAN I DO WITH OPERATOR SDK?

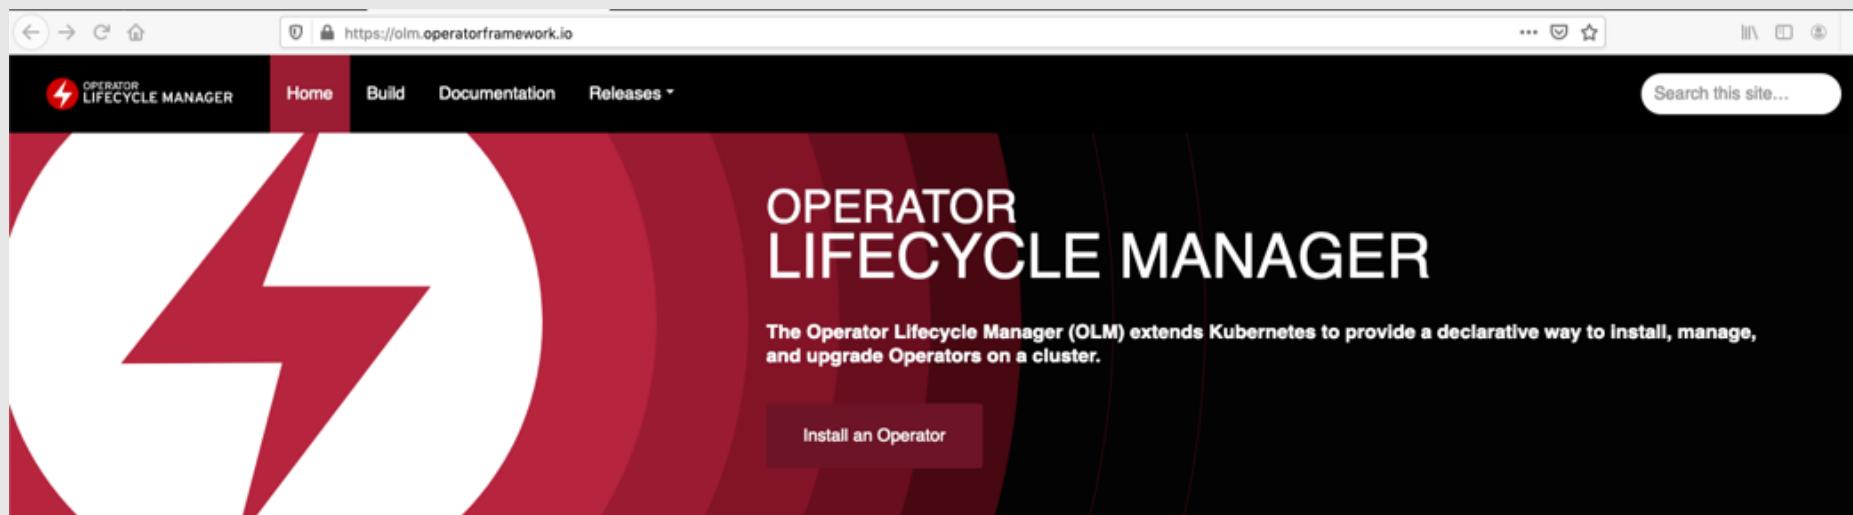
The Operator SDK provides the tools to build, test, and package Operators. Initially, the SDK facilitates the marriage of an application's business logic (for example, how to scale, upgrade, or backup) with the Kubernetes API to execute those operations. Over time, the SDK can allow engineers to make applications smarter and have the user experience of cloud services. Leading practices and code patterns that are shared across Operators are included in the SDK to help prevent reinventing the wheel.

The Operator SDK is a framework that uses the controller-runtime library to make writing operators easier by providing:

- High level APIs and abstractions to write the operational logic more intuitively
- Tools for scaffolding and code generation to bootstrap a new project fast
- Extensions to cover common Operator use cases



Operator Lifecycle Manager (OLM)



The screenshot shows the official website for Operator Lifecycle Manager (OLM) at <https://olm.operatorframework.io>. The page features a large red lightning bolt icon on the left. The main title "OPERATOR LIFECYCLE MANAGER" is prominently displayed in white. Below it, a subtitle reads: "The Operator Lifecycle Manager (OLM) extends Kubernetes to provide a declarative way to install, manage, and upgrade Operators on a cluster." A "Install an Operator" button is visible. The top navigation bar includes links for Home, Build, Documentation, and Releases.

WHAT IS OPERATOR LIFECYCLE MANAGER?

This project is a component of the [Operator Framework](#), an open source toolkit to manage Kubernetes native applications, called Operators, in a streamlined and scalable way.

OLM FEATURES

OVER-THE-AIR UPDATES AND CATALOGS	DEPENDENCY MODEL	DISCOVERABILITY	CLUSTER STABILITY	DECLARATIVE UI CONTROLS
OLM provides rich update mechanisms to keep Kubernetes native applications up to date automatically.	With OLM's packaging format Operators can express dependencies on the platform and on other Operators.	OLM makes Operators and their services available for cluster users to select and install.	OLM will prevent conflicting Operators owning the same APIs being installed, ensuring cluster stability.	OLM enables Operators to behave like managed service providers through the APIs they expose.



OperatorHub.io

https://operatorhub.io

Welcome to OperatorHub.io

OperatorHub.io is a new home for the Kubernetes community to share Operators. Find an existing Operator or list your own today.

CATEGORIES 179 ITEMS VIEW SORT A-Z

AI/Machine Learning Application Runtime Big Data Cloud Provider Database Developer Tools Integration & Delivery Logging & Tracing Monitoring Networking OpenShift Optional Security Storage Streaming & Messaging

PROVIDER

- Alibaba Cloud (1)
- Altinity (1)
- Anchore (1)
- Apicurio (1)
- AppDynamics (1)

Show 115 more

[DEPRECATED] Knative Eventing Operator provided by The Knative Authors	[DEPRECATED] Knative Serving Operator provided by The Knative Authors	Akka Cluster Operator provided by Lightbend, Inc.	Altinity ClickHouse Operator provided by Altinity	Anchore Engine Operator provided by Anchore Inc.
Apache Spark Operator provided by radanalytics.io	API Operator for Kubernetes provided by WSO2	APIcast provided by Red Hat	Apicurio Registry Operator provided by Apicurio	Appdynamics Operator provided by AppDynamics LLC
An operator for managing the Apache Spark clusters and intelligent applications th...	API Operator provides a fully automated experience for...	APIcast is an API gateway built on top of NGINX. It is part of the 3scale API Management...	Deploy and manage Apicurio Registry on Kubernetes.	End to end monitoring of applications on Kubernetes and OpenShift clusters with...
Appranix CPS Operator provided by Appranix, Inc	Appsody Operator provided by Appsody	Aqua Security Operator provided by Aqua Security, Inc.	ArangoDB provided by ArangoDB GmbH	Argo CD provided by Argo CD



Alternative Frameworks

- KOPF - Kubernetes Operator Pythonic Framework
 - Implement in Python in operator.py file
 - Must manually create all boilerplate files - crd, role, rolebinding, etc.
- Java Operator SDK
 - Does some scaffolding
 - Implement in Java
 - Still fairly new
- Java or other language
 - Requires writing more code
 - Have to implement all files
 - Single programming language
- Also Kubebuilder, but now incorporated in SDK



Operator SDK vs kubebuilder

- Formerly separate approaches for building operators
- kubebuilder
 - a kubernetes-sig (part of API Machinery SIG)
 - out of Google
- operator-sdk
 - a part of the operator-framework - CNCF (Cloud Native Computing Foundation) project
 - originated by CoreOS and incorporated/managed now by Red Hat
- SDK has adopted kubebuilder approach and mechanisms
- SDK is “consumer” of kubebuilder
- SDK refers to kubebuilder documentation
- SDK maintainers contribute back to kubebuilder code
- Beyond kubebuilder functionality, SDK contains functionality for Operator Lifecycle Management





General Operator-SDK (kubebuilder) flow

- Prereqs - Go, Docker, kubectl, cluster access
- Install operator-sdk
- Create a project
- Create an API
- Test it out
- Install instances of Custom Resources
- Run it on the Cluster



Operator-SDK Languages - Go

- Go - general-purpose programming
 - can code any logic
 - good match for Kubernetes api client library since Kubernetes is written in Go
 - scaffold Go project
 - » `$ operator-sdk init --repo=<Git repo>`
 - » `$ operator-sdk create api --group <name> --version <version> --kind <kind> --resource=true --controller=true`
 - Add logic for operator into Reconcile function in controllers/<name>_controller.go
 - Add default parameters into api/v1/<name>_types.go
- Notes:
 - Imperative approach - more work required and more complicated
 - Controller code is scaffolded but not sufficient
 - Go is well integrated with K8s
 - Essentially can do any logic without limits
 - Operator configuration files (role, role-binding, crd, etc.) are automatically generated



Workflow

- ❑ Use sdk to create scaffold project (\$ operator-sdk init)
- ❑ Create scaffold API , CRD, and controller (\$ operator-sdk create api)
- ❑ Update API to have necessary fields for spec and status (\$ edit *_types.go)
- ❑ After above, automatically update generated code for resource - updates interface/implementation for runtime.Object to be a valid Kubernetes Kind (\$ make generate)
 - ❑ updates api/<version>/zz_generated.deepcopy.go
- ❑ Generate CRD manifests (\$ make manifests)
 - ❑ creates config/crd/bases/<item>.yaml
 - ❑ creates config/rbac/role.yaml
- ❑ Update scaffolded controller code to complete controller functionality
 - ❑ Add needed packages and core functions from Kubernetes
 - ❑ Add code to reconcile desired state of CR against observed state
 - ❑ Add code to create any needed K8s objects (Kinds)
 - ❑ Add any additional markers to manage items such as permissions/generate manifests
 - ❑ updates controllers/<item>_controller.go
- ❑ Update CRD manifests (\$ make manifests)
 - ❑ updates config/crd/bases/<item>.yaml
 - ❑ updates config/rbac/role.yaml
- ❑ Execute install target to register CRD with K8s apiserver (\$ make install)
- ❑ Test the operator (\$ make run)
- ❑ Create the docker package (\$ make docker-build)
- ❑ Push image to registry (\$ make docker-push)
- ❑ Deploy operator to cluster (\$ make deploy)

Lab 1 – Scaffolding an Operator with the Operator SDK



Comparison of Languages for Operators using Operator-SDK

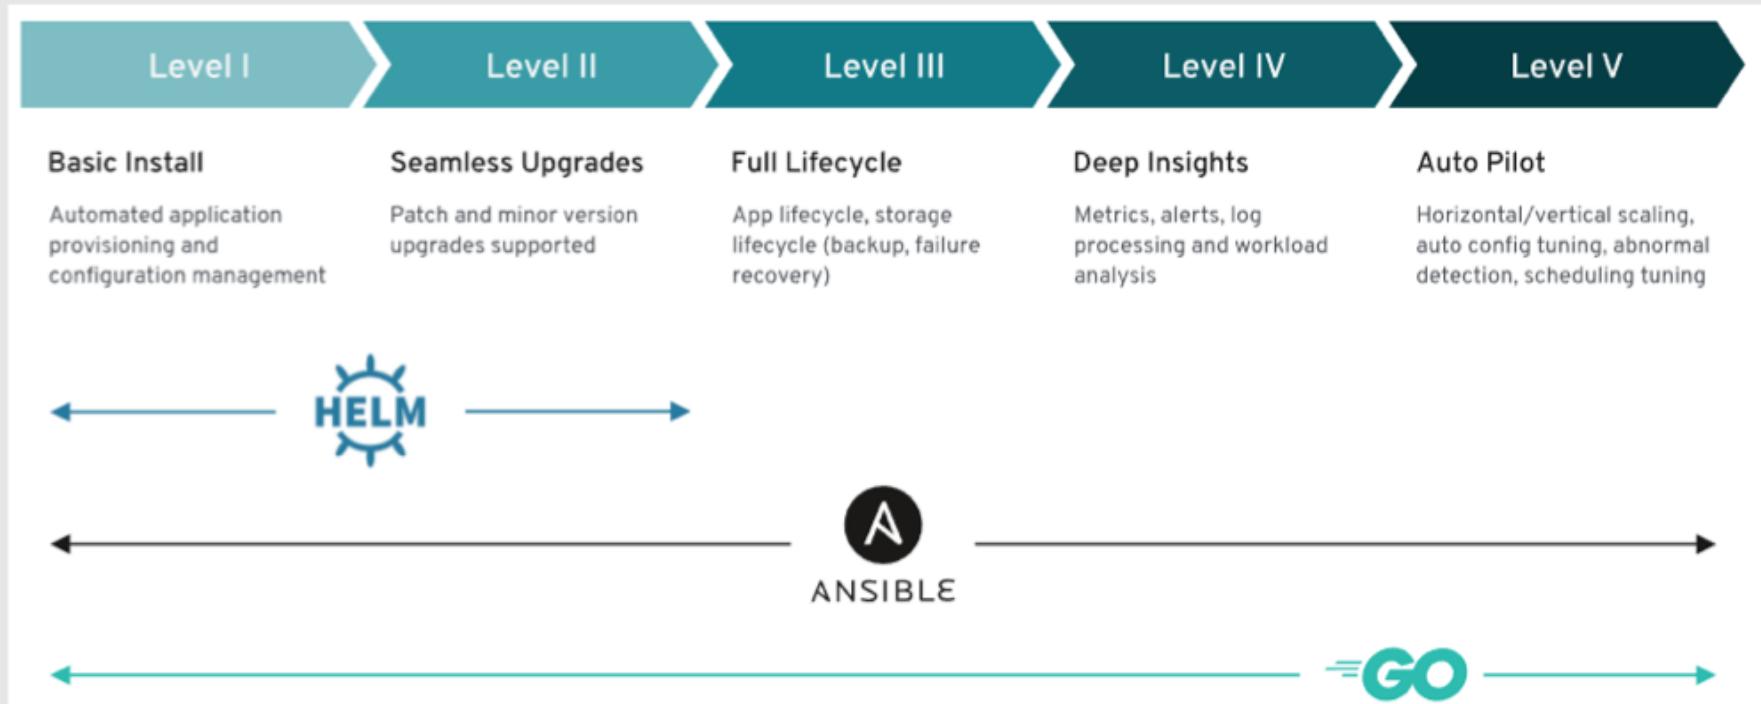
- source: operatorhub.io

Operator Type	What the SDK generates	What you need to define
Go Operator	<ul style="list-style-type: none">General go program structureBoilerplate code to talk to the Kubernetes APIBoilerplate code to watch for Kubernetes objects of interestAn entry point to the reconciliation loopAn example YAML files based on CRDs	<ul style="list-style-type: none">Custom objects via CRDsControl loop logic in GoPotentially artistic stunts only possible by talking directly to the API from Go
Ansible Operator	<ul style="list-style-type: none">A Go program that runs an Ansible playbook or role every time a certain type of object is detected / modified	<ul style="list-style-type: none">Ansible playbook or roleCustom objects via CRD
Helm Operator	<ul style="list-style-type: none">A custom object via CRD containing the same properties as the chart's values.yamlA Go program that reads a helm chart and deploys all its resourcesWatch statements to detect changes in the custom objects specification, re-deploying all resources with updated values	<ul style="list-style-type: none">The location / repository of the helm chart



Operator Maturity Model

- source: operatorhub.io





Parts of the project

- Build infrastructure
 - go.mod - Go module for project with basic dependencies
 - Makefile - make targets for building and deploying the controller
 - PROJECT - metadata for new component scaffolds
- Launch configuration
 - Initially, just has Kustomize YAML to launch controller on cluster
 - After starting to write controller, has
 - » Custom Resource Definitions
 - » RBAC config
 - » Webhook config
 - config/default - Kustomize base for launching controller (standard config)
 - config/manager - Kustomize base to launch controller as pods in cluster
 - config/rbac - Kustomize base for service account to run controllers as
- Entrypoint - main.go



Initializing a project w/ operator-sdk init

- **bin/manager** - executable to wrap 1 or more controllers
- **config/** - launch configurations
 - initially just contains Kustomize YAML definitions needed to launch controller in cluster
 - eventually will also contain CRD, RBAC configs, WebhookConfigs
- **config/certmanager** - files for K8s add-on (cert-manager) that automates management and issuance of TLS certs from various sources
- **config/default** - Kustomize base to launch controller in standard config
- **config/manager** - launch controllers as pods in cluster
- **config/prometheus** - patch file to enable prometheus metrics via ServiceMonitor
- **config/rbac** - permissions req'd to run controllers under their own SA
- **config/webhook** - for deploying admission webhooks
- **Dockerfile** - for building the manager binary
- **go.mod** - project dependencies
- **go.sum** - cryptographic checksums of specific module versions
- **hack/boilerplate.go.txt** - Apache 2 license file example
- **main.go** - basic entrypoint for project
- **Makefile** - Make targets to build and deploy controller and more
- **PROJECT** - Kubebuilder metadata for scaffolding new components

```
diyuser3@training1:~/ops$ tree
.
├── bin
│   └── manager
├── config
│   ├── certmanager
│   │   ├── certificate.yaml
│   │   ├── kustomization.yaml
│   │   └── kustomizeconfig.yaml
│   ├── default
│   │   ├── kustomization.yaml
│   │   ├── manager_auth_proxy_patch.yaml
│   │   ├── manager_webhook_patch.yaml
│   │   └── webhookcainjection_patch.yaml
│   ├── manager
│   │   ├── kustomization.yaml
│   │   └── manager.yaml
│   ├── prometheus
│   │   ├── kustomization.yaml
│   │   └── monitor.yaml
│   └── rbac
│       ├── auth_proxy_client_clusterrole.yaml
│       ├── auth_proxy_role_binding.yaml
│       ├── auth_proxy_role.yaml
│       ├── auth_proxy_service.yaml
│       ├── kustomization.yaml
│       ├── leader_election_role_binding.yaml
│       ├── leader_election_role.yaml
│       └── role_binding.yaml
└── webhook
    ├── kustomization.yaml
    ├── kustomizeconfig.yaml
    └── service.yaml
Dockerfile
go.mod
go.sum
hack
└── boilerplate.go.txt
main.go
Makefile
PROJECT
```



Groups and Versions and Kinds

- <https://book.kubebuilder.io/cronjob-tutorial/gvks.html>
- API Group in K8s
 - collection of related functionality
 - one or more versions - denotes version of API over time
- Kinds
 - API types contained in group-version
 - Must be backwards compatible with previous API versions
- Resource
 - Use of a Kind in the API
- GVK
 - GroupVersionKind
 - corresponds to given root Go type in package
- Scheme
 - tracks Go type to GVK



Spec and Status

- Every running k8s object includes two nested object fields
 - spec
 - User-provided
 - Describes desired state for the object
 - status
 - Provided and updated by k8s system
 - Describes the actual state of the object

```
diyuser3@training1:~$ kubectl edit -n roar2 pod roar-web-74bb47bdb8-56l4n
```

```
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: "2019-07-21T20:04:34Z"
  generateName: roar-web-74bb47bdb8-
  labels:
    app: roar-web
    pod-template-hash: 74bb47bdb8
  name: roar-web-74bb47bdb8-56l4n
  namespace: roar2
  ownerReferences:
    - apiVersion: apps/v1
      blockOwnerDeletion: true
      controller: true
      kind: ReplicaSet
      name: roar-web-74bb47bdb8
      uid: 89d00855-abf0-11e9-b30c-080027c4188a
  resourceVersion: "346085"
  selfLink: /api/v1/namespaces/roar2/pods/roar-web-74bb47bdb8-56l4n
  uid: c2b50783-abf2-11e9-b30c-080027c4188a
spec:
  containers:
    - image: localhost:5000/roar-web:v2
      imagePullPolicy: Always
      name: roar-web
      ports:
        - containerPort: 8080
          name: web
          protocol: TCP
status:
  secret:
    defaultMode: 420
    secretName: default-token-dxx7t
  status:
    conditions:
      - lastProbeTime: null
        lastTransitionTime: "2019-07-21T20:04:34Z"
        status: "True"
        type: Initialized
      - lastProbeTime: null
        lastTransitionTime: "2019-07-21T20:04:34Z"
        status: "True"
        type: Ready
      - lastProbeTime: null
        lastTransitionTime: "2019-07-21T20:04:34Z"
        status: "True"
        type: ContainersReady
      - lastProbeTime: null
        lastTransitionTime: "2019-07-21T20:04:34Z"
        status: "True"
        type: PodScheduled
    containerStatuses:
      - containerID: docker://41c2d31ef46b8ea1
        image: localhost:5000/roar-web:v2
        imageID: docker-pullable://localhost:5000/roar-web:v2
        lastState:
          terminated:
            containerID: docker://c0ddd9686971
            exitCode: 143
```



operator-sdk create api

- scaffolds new Kind and related controller
- expects group, version, and kind
- prompts to create resource & controller
- in api/v1alpha1
 - new file for Kind - roarapp_types.go
 - groupversion_info.go - common metadata about group-version
 - zz_generated.deepcopy.go
 - » contains automatically generated code of runtime.Object interface
 - » needed for K8s to recognize as a Kind
- in controllers
 - roarapp_controller.go - scaffolded source for controller
 - suite_test.go - does the bare bones of setting up a test environment

```
diyuser3@training1:~/ops$ tree
.
└── api
    └── v1alpha1
        ├── groupversion_info.go
        ├── roarapp_types.go
        └── zz_generated.deepcopy.go
```

```
└── controllers
    └── roarapp_controller.go
        └── suite_test.go
```



What's in an API? (_types.go)

- License & imports
- spec
 - holds desired state
 - “inputs” to controller
 - for our app
 - » number of replicas
 - » image for webapp
 - » image for backend db
- status
 - holds observed state
 - info that we want users or other controllers to have access to
- schema
 - defines type and list of type
 - TypeMeta - describes API version & Kind
 - ObjectMeta - holds items like name, namespace, labels
- init
 - adds Go types to API group
 - allows adding types in this group to any schema

```
// RoarAppSpec defines the desired state of RoarApp
type RoarAppSpec struct {
    // INSERT ADDITIONAL SPEC FIELDS - desired state of cluster
    // Important: Run "make" to regenerate code after modifying this file
    Replicas int32 `json:"replicas"`
    WebImage string `json:"webImage,omitempty"`
    DbImage  string `json:"dbImage,omitempty"`
}

// RoarAppStatus defines the observed state of RoarApp
type RoarAppStatus struct {
    // INSERT ADDITIONAL STATUS FIELD - define observed state of cluster
    // Important: Run "make" to regenerate code after modifying this file
    PodNames []string `json:"podNames"`
}

// RoarApp is the Schema for the roarapps API
type RoarApp struct {
    metav1.TypeMeta `json:",inline"`
    metav1.ObjectMeta `json:"metadata,omitempty"`

    Spec RoarAppSpec `json:"spec,omitempty"`
    Status RoarAppStatus `json:"status,omitempty"`
}

// +kubebuilder:object:root=true
// RoarAppList contains a list of RoarApp
type RoarAppList struct {
    metav1.TypeMeta `json:",inline"`
    metav1.ListMeta `json:"metadata,omitempty"`
    Items          []RoarApp `json:"items"`
}

// +kubebuilder - marker - tells object generator that this type is a Kind
func init() {
    SchemeBuilder.Register(&RoarApp{}, &RoarAppList{})
}
```

Scheme - methods for serializing, converts to/from Go schemas



Custom Resource

- Background
 - In k8s, a resource is an endpoint in the API that stores groups of API objects of a specific kind
 - Objects can be managed by kubectl
- A new kind of object (other than standard ones such as Pod, Service, etc.) targeted around a resource you want to run and manage in Kubernetes
- Typically requires special handling to respond to Kubernetes operations (may be stateful) such as scaling, deleting, etc.
- Actually 2 methods to add these -
 - Custom Resource Definitions (CRDs) - easier to implement/share
 - Aggregated API - more flexible, requires programming



Makefile

- Convenient way to wrap operations
- Used for SDK Go process
- Leverages tools like controller-gen and kustomize

```
# Current Operator version
VERSION ?= 0.0.1
# Default bundle image tag
BUNDLE_IMG ?= controller-bundle:${VERSION}
# Options for 'bundle-build'
ifeq ($(origin CHANNELS), undefined)
BUNDLE_CHANNELS := --channels=${CHANNELS}
endif
ifeq ($(origin DEFAULT_CHANNEL), undefined)
BUNDLE_DEFAULT_CHANNEL := --default-channel=${DEFAULT_CHANNEL}
endif
BUNDLE_METADATA_OPTS ?= ${BUNDLE_CHANNELS} ${BUNDLE_DEFAULT_CHANNEL}

# Image URL to use all building/pushing image targets
IMG ?= controller:latest
# Produce CRDs that work back to Kubernetes 1.11 (no version conversion)
CRD_OPTIONS ?= "crd:trivialVersions=true"

# Get the currently used golang install path (in GOPATH/bin, unless GOBIN is set)
ifeq (,$(shell go env GOBIN))
GOBIN=$(shell go env GOPATH)/bin
else
GOBIN=$(shell go env GOBIN)
endif

all: manager

# Run tests
test: generate fmt vet manifests
    go test ./... -coverprofile cover.out

# Build manager binary
manager: generate fmt vet
    go build -o bin/manager main.go

# Run against the configured Kubernetes cluster in ~/.kube/config
run: generate fmt vet manifests
    go run ./main.go

# Install CRDs into a cluster
install: manifests kustomize
    $(KUSTOMIZE) build config/crd | kubectl apply -f -

# Uninstall CRDs from a cluster
uninstall: manifests kustomize
    $(KUSTOMIZE) build config/crd | kubectl delete -f -

# Deploy controller in the configured Kubernetes cluster in ~/.kube/config
deploy: manifests kustomize
    cd config/manager && $(KUSTOMIZE) edit set image controller=${IMG}
    $(KUSTOMIZE) build config/default | kubectl apply -f -

# Generate manifests e.g. CRD, RBAC etc.
manifests: controller-gen
    $(CONTROLLER_GEN) ${CRD_OPTIONS} rbac:roleName=manager-role webhookPaths="./..."
```



Workflow

- ✓ Use sdk to create scaffold project (\$ operator-sdk init)
- ✓ Create scaffold API , CRD, and controller (\$ operator-sdk create api)
- ❑ Update API to have necessary fields for spec and status (\$ edit *_types.go)
- ❑ After above, automatically update generated code for resource - updates interface/implementation for runtime.Object to be a valid Kubernetes Kind (\$ make generate)
 - ❑ updates api/<version>/zz_generated.deepcopy.go
- ❑ Generate CRD manifests (\$ make manifests)
 - ❑ creates config/crd/bases/<item>.yaml
 - ❑ creates config/rbac/role.yaml
- ❑ Update scaffolded controller code to complete controller functionality
 - ❑ Add needed packages and core functions from Kubernetes
 - ❑ Add code to reconcile desired state of CR against observed state
 - ❑ Add code to create any needed K8s objects (Kinds)
 - ❑ Add any additional markers to manage items such as permissions/generate manifests
 - ❑ updates controllers/<item>_controller.go
- ❑ Update CRD manifests (\$ make manifests)
 - ❑ updates config/crd/bases/<item>.yaml
 - ❑ updates config/rbac/role.yaml
- ❑ Execute install target to register CRD with K8s apiserver (\$ make install)
- ❑ Test the operator (\$ make run)
- ❑ Create the docker package (\$ make docker-build)
- ❑ Push image to registry (\$ make docker-push)
- ❑ Deploy operator to cluster (\$ make deploy)

Lab 2 - Defining the API



What happens when we create a CRD?

- Create by defining spec and applying it
- CRD can be either scoped for a namespace or cluster
 - » Per “scope” field
- K8S api server creates a new RESTful resource path for each version you supply
- Creates new RESTful API endpoint
 - » Example:

/apis/stable.example.com/v1/namespaces/*/crontab
s/...

- Endpoint URL can then be used to create and manage custom objects
- Manipulate as other objects
- Example Kind objects will be CronTab
- Can then make instances of these custom objects
- And use them with operations like kubectl get
- Creating these made easier with SDK tools, such as “controller-runtime” & “controller-gen”

```
apiVersion: "stable.example.com/v1"
kind: CronTab
metadata:
  name: my-new-cron-object
spec:
  cronSpec: "* * * * */5"
  image: my-awesome-cron-image
```

```
1  apiVersion: apiextensions.k8s.io/v1
2  kind: CustomResourceDefinition
3  metadata:
4    # name must match the spec fields below, and be in the form: <group>.stable.example.com
5  name: crontabs.stable.example.com
6  spec:
7    # group name to use for REST API: /apis/<group>/<version>
8    group: stable.example.com
9    # list of versions supported by this CustomResourceDefinition
10   versions:
11     - name: v1
12       # Each version can be enabled/disabled by Served flag.
13       served: true
14       # One and only one version must be marked as the storage version
15       storage: true
16       schema:
17         openAPIV3Schema:
18           type: object
19           properties:
20             spec:
21               type: object
22               properties:
23                 cronSpec:
24                   type: string
25                 image:
26                     type: string
27                 replicas:
28                   type: integer
29       # either Namespaced or Cluster
30       scope: Namespaced
31       names:
32         # plural name to be used in the URL: /apis/<group>/<version>/<plural>
33         plural: crontabs
34         # singular name to be used as an alias on the CLI and for display
35         singular: crontab
36         # kind is normally the CamelCased singular type. Your resource
37         kind: CronTab
38         # shortNames allow shorter string to match your resource on the CLI
39         shortNames:
40           - ct
```



What is the controller-runtime?

- set of Go libraries for building Controllers
- used by operator-sdk and kubebuilder
- referenced via the Controller Runtime Client API
- provides ways to watch and reconcile resources in a K8s cluster
- supports CRUD (Create, [Get, List], Update, Delete) operations
- operators use at least 1 controller to perform task
 - generally through some combination of CRUD operations
- Operator SDK uses controller-runtime's Client interface
 - provides interface for these operations



What is controller-gen?

- Tool (from KubeBuilder) for generating K8s YAML and utility code
- Part of larger set of Go libraries for generating controllers - **controller-tools**
- Code and config generation controlled by “marker comments” in Go code
- Types of items produced via controller-gen
 - operator role
 - CRD
 - add validation to CRD
 - scaffold DeepCopy functions for objects
 - help in producing webhooks
- Controller-gen is composed of:
 - Generators - specify what to generate
 - Output rules - specify where and how to write the results
- Integration with make in sdk
 - make manifests - generates K8s object YAML
 - » CRDs, webhook configurations, RBAC roles
 - make generate - generates code
 - » runtime.Object/DeepCopy implementations



What makes up an operator controller?

- Standard imports

- core controller-runtime library
- client package
- package for API types

```
import (
    "context"

    "github.com/go-logr/logr"
    "k8s.io/apimachinery/pkg/runtime"
    ctrl "sigs.k8s.io/controller-runtime"
    "sigs.k8s.io/controller-runtime/pkg/client"

    roarappv1alpha1 "github.com/brentlaster/roar-op/api/v1alpha1"
)
```

- Reconciler data structure

- Scaffolded to start
- Includes logging and fetching objs
- Schemes - associate Go types with K8s groups, versions, and kinds

```
// RoarAppReconciler reconciles a RoarApp object
type RoarAppReconciler struct {
    client.Client
    Log    logr.Logger
    Scheme *runtime.Scheme
}
```

- RBAC permissions

- via RBAC markers
- Starts with bare minimum

```
// +kubebuilder:rbac:groups=roarapp.roarapp.com,resources=roarapps,verbs=get;list;watch;create;update;patch;delete
// +kubebuilder:rbac:groups=roarapp.roarapp.com,resources=roarapps/status,verbs=get;update;patch
```

- Reconciler function

- Request contains info needed to uniquely identify an object - Name and Namespace
- Context allows cancellation of requests, tracing, etc.
 - » Background context is basic one
- logging handle uses structured logging library named "logr"
- logging works via attaching key-value pairs to static message

```
func (r *RoarAppReconciler) Reconcile(req ctrl.Request) (ctrl.Result, error) {
    _ = context.Background()
    _ = r.Log.WithValues("roarapp", req.NamespacedName)

    // your logic here

    return ctrl.Result{}, nil
}
```

```
func (r *RoarAppReconciler) SetupWithManager(mgr ctrl.Manager) error {
    return ctrl.NewControllerManagedBy(mgr).
        For(&roarappv1alpha1.RoarApp{}).
        Complete(r)
}
```

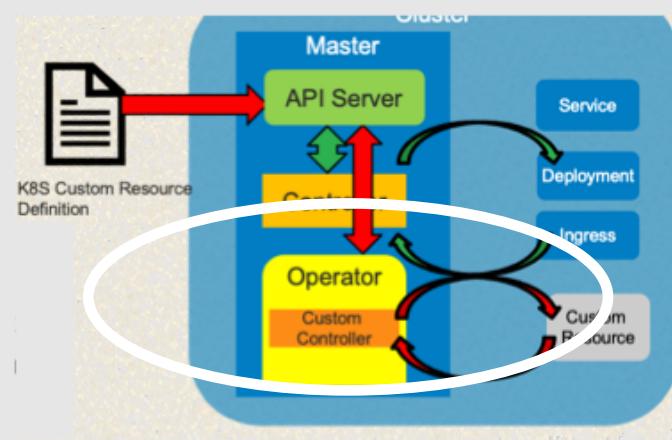
- SetupWithManager function

- Adds reconciler to manager
- Want it started when manager is started



Reconcile and Client API

- Reconciler implements reconcile.Reconciler interface
 - exposes Reconcile method
 - added to corresponding Controller for a Kind
 - called in response to cluster or external Events
 - reads and writes cluster state by Controller
 - returns ctrl.Result
 - SDK Reconcilers have Client access to make K8s API calls



```
→ C ⌂ https://github.com/brentlaster/ops-ref/blob/main/controllers/roarapp_controller.go  
57  
58 func (r *RoarAppReconciler) Reconcile(req ctrl.Request) (ctrl.Result, error) {  
59  
60     log := r.Log.WithValues("roarapp", req.NamespacedName)  
61  
62     log.Info("Reconciling instance")  
63  
119     if numAvailable > instance.Spec.Replicas {  
120         log.Info("Scaling down pods", "Currently available", numAvailable, "Required replicas", instance.Spec.Replicas)  
121         diff := numAvailable - instance.Spec.Replicas  
122         dpods := available[:diff]  
151     if numAvailable < instance.Spec.Replicas {  
152         log.Info("Scaling up pods", "Currently available", numAvailable, "Required replicas", instance.Spec.Replicas)  
153         // Define a new Pod object  
154         pod := newPodForCR(instance)
```



Operators vs Custom Controllers

- Operators are controllers
 - All operators use the controller pattern
 - Not all controllers are operators
- Operators codify operational knowledge for an application
- Operators
 - Encapsulate application-specific knowledge
 - Manage lifecycle for an application
 - Provide a CRD
- Operator = controller pattern + API extension + single-app focus
- Controllers work on existing K8s resources, operators on custom
- Custom controller - program that subscribes and listens to API server
 - When detects event it is interested in, takes action, may use information found in event's data



Workflow

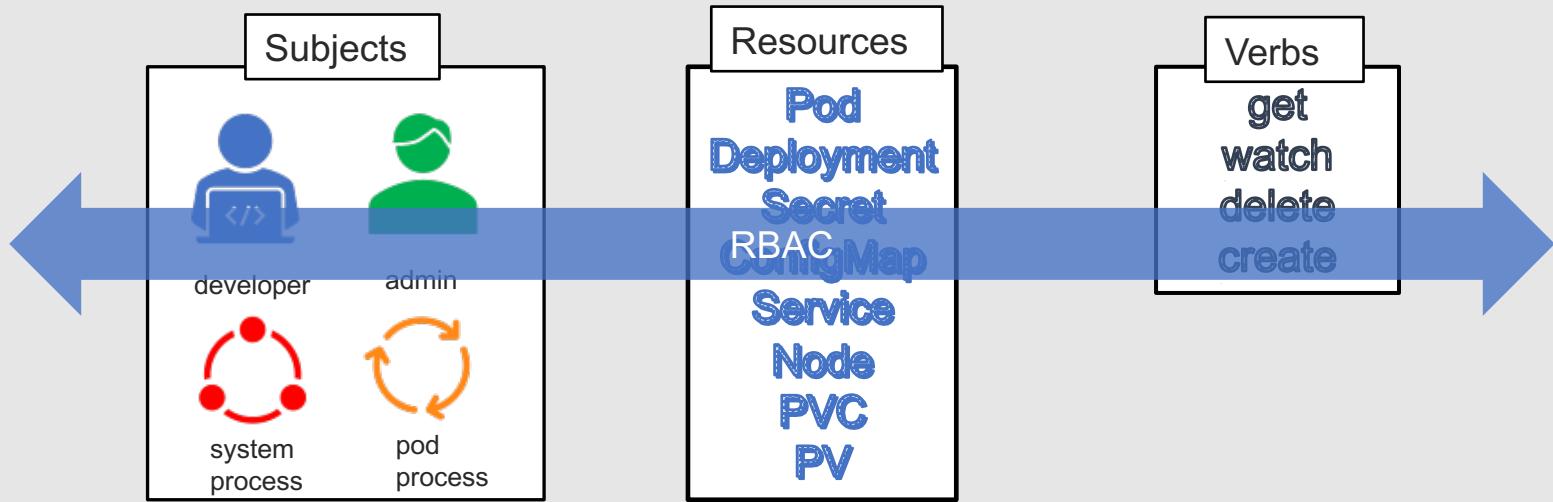
- ✓ Use sdk to create scaffold project (\$ operator-sdk init)
- ✓ Create scaffold API , CRD, and controller (\$ operator-sdk create api)
- ✓ Update API to have necessary fields for spec and status (\$ edit *_types.go)
- ✓ After above, automatically update generated code for resource - updates interface/implementation for runtime.Object to be a valid Kubernetes Kind (\$ make generate)
 - ✓ updates api/<version>/zz_generated.deepcopy.go
- ✓ Generate CRD manifests (\$ make manifests)
 - ✓ creates config/crd/bases/<item>.yaml
 - ✓ creates config/rbac/role.yaml
- ❑ Update scaffolded controller code to complete controller functionality
 - ❑ Add needed packages and core functions from Kubernetes
 - ❑ Add code to reconcile desired state of CR against observed state
 - ❑ Add code to create any needed K8s objects (Kinds)
 - ❑ Add any additional markers to manage items such as permissions/generate manifests
 - ❑ updates controllers/<item>_controller.go
- ❑ Update CRD manifests (\$ make manifests)
 - ❑ updates config/crd/bases/<item>.yaml
 - ❑ updates config/rbac/role.yaml
- ❑ Execute install target to register CRD with K8s apiserver (\$ make install)
- ❑ Test the operator (\$ make run)
- ❑ Create the docker package (\$ make docker-build)
- ❑ Push image to registry (\$ make docker-push)
- ❑ Deploy operator to cluster (\$ make deploy)

Lab 3 - Implementing the Controller



Kubernetes RBAC

- The pieces
 - 3 types of elements involved
 - Subjects: users or processes that need to access the API
 - Resources: Kubernetes API Objects available in the cluster. (Examples include Pods, Deployments, Services, Nodes, and Persistent Volumes, etc.)
 - Verbs: The Kubernetes operations that can be executed on the resources above. Multiple kinds of verbs available – all are CRUD operations (Examples include get, watch, create, delete, etc.)

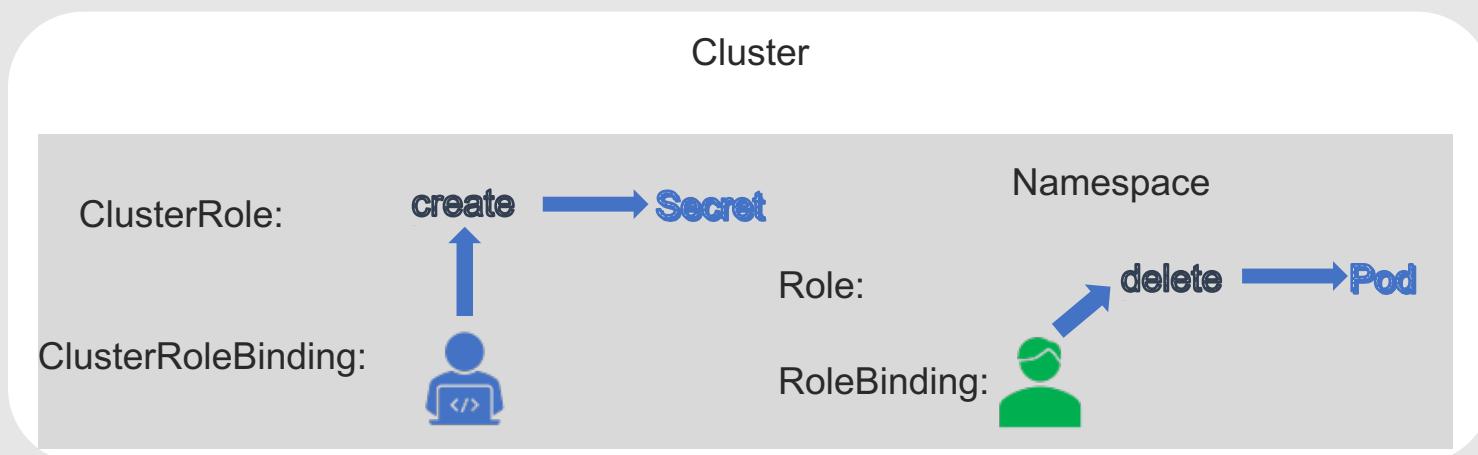


Goal: Connect subjects, API resources, and verbs. Specify, given a user, which operations can be executed over which resources.



Kubernetes RBAC

- The mechanisms
 - Object types
 - k8s Roles: Connect API Resources and Verbs; can be reused for different Subjects; bound to a namespace
 - k8s RoleBinding: Connects Subjects to Roles; Given a role that combines API objects and verbs, defines which subjects can use it.
 - k8s ClusterRole: a role to be applied across a cluster
 - k8s ClusterRoleBinding: Connects Subjects to ClusterRole.





Kubernetes RBAC - Service Accounts

- Users: These are global, and meant for humans or processes living outside the cluster.
- ServiceAccounts: These are namespaced and meant for intra-cluster processes running inside pods.
- Both authenticate against the API to get access to resources
- But there is no kubernetes “User” object
- So you can do “kubectl create serviceaccount <name>” but not “kubectl create user <name>”
- Bottom line: cluster does not store any information about users, so must be managed outside of cluster with certificates, tokens, etc.

Create ServiceAccount

```
apiVersion: v1
kind: ServiceAccount
metadata:
  name: myapp
  labels:
    app: myapp
```

Create Role

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: myapp-access-ep
  labels:
    app: myapp
rules:
- apiGroups: []
  resources: ["endpoints"]
  verbs: ["get"]
```

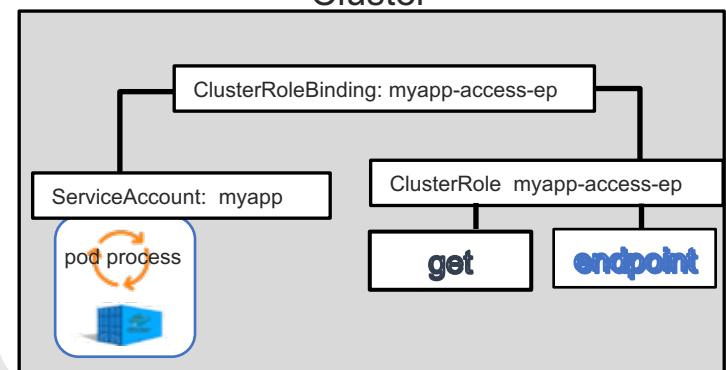
Add ServiceAccount to Pod

```
[...]
spec:
  template:
    metadata:
      name: myapp
      labels:
        app: myapp
    spec:
      serviceAccountName: myapp
      containers:
[...]
```

Tie Role to ServiceAccount

```
kind: ClusterRoleBinding
apiVersion: rbac.authorization.k8s.io/v1
metadata:
  name: myapp-access-ep
  labels:
    app: myapp
subjects:
- kind: ServiceAccount
  name: myapp
roleRef:
  kind: ClusterRole
  name: myapp-access-ep
  apiGroup: rbac.authorization.k8s.io
```

Cluster





Markers for Config/Code Generation

- Carried over in sdk from Kubebuilder
- Markers are special comments in Go code
 - start with “// +”
 - indicate additional info around types, fields, and packages
- Signal “controller-gen” tool to generate
 - utility code
 - » done via “make generate”
 - » examples: runtime.Object/DeepCopy implementations
 - object YAML
 - » done via “make manifests”
 - » examples: CRDs, RBAC roles



The screenshot shows a browser window with the URL https://github.com/brentlaster/ops-ref/blob/main/controllers/roarapp_controller.go. The code snippet displays several lines of Go code with markers. Lines 54, 55, and 56 show markers for RBAC permissions. Line 58 starts a function definition.

```
54 // +kubebuilder:rbac:groups=core,resources=pods,verbs=get;list;watch;create;update;delete
55 // +kubebuilder:rbac:groups=batch,resources=jobs,verbs=get;list;watch;create;update;patch;delete
56 // +kubebuilder:rbac:groups=core,resources=services,verbs=get;list;watch;create;update;delete
57
58 func (r *RoarAppReconciler) Reconcile(req ctrl.Request) (ctrl.Result, error) {
59 }
```



Marker Syntax

- 3 types
 - Empty (+kubebuilder:validation:Optional)
 - » behave like boolean flags - presence indicates "on"
 - Anonymous (+kubebuilder:validation:Maximum=20)
 - » take a single value as argument
 - Multi-option (+kubebuilder:printcolumn:JSONPath=".status.replicas",name=Replicas,type=string)
 - » take one or more arguments
 - » first argument separated from name by colon
 - » remaining arguments separated by comma
 - » order doesn't matter

```
43 // +kubebuilder:subresource:status
44
45 // RoarApp is the Schema for the roarapps API
46 type RoarApp struct {
47     metav1.TypeMeta `json:",inline"`
48     metav1.ObjectMeta `json:"metadata,omitempty"`
49
50     Spec   RoarAppSpec   `json:"spec,omitempty"`
51     Status RoarAppStatus `json:"status,omitempty"`
52 }
```



```
54 // +kubebuilder:object:root=true
55
56 // RoarAppList contains a list of RoarApp
57 type RoarAppList struct {
58     metav1.TypeMeta `json:",inline"`
59     metav1.ListMeta `json:"metadata,omitempty"`
60     Items           []RoarApp `json:"items"`
61 }
```



```
54 // +kubebuilder:rbac:groups=core,resources=pods,verbs=get;list;watch;create;update;delete
55 // +kubebuilder:rbac:groups=batch,resources=jobs,verbs=get;list;watch;create;update;patch;delete
56 // +kubebuilder:rbac:groups=core,resources=services,verbs=get;list;watch;create;update;delete
57
58 func (r *RoarAppReconciler) Reconcile(req ctrl.Request) (ctrl.Result, error) {
59 }
```



Marker Generator Categories

CRD Generation

These markers describe how to construct a custom resource definition from a series of Go types and packages. Generation of the actual validation schema is described by the [validation markers](#).

See [Generating CRDs](#) for examples.

- ▶ Show Detailed Argument Help

```
// +kubebuilder:printcolumn  
:JSONPath=<string>,description=<string>,format=<string>,name=<string>,priority=<int>,type=<string>  
on type  
    adds a column to "kubectl get" output for  
    this CRD.
```

```
// +kubebuilder:resource  
:categories=</>string>,path=<string>,scope=<string>  
,shortName=</>string>,singular=<string>  
on type  
    configures naming and scope for a CRD.
```

```
// +kubebuilder:skipversion on type  
    ▶ removes the particular version of the CRD  
    from the CRDs spec.
```

```
// +kubebuilder:storageversion on type  
    ▶ marks this version as the "storage  
    version" for the CRD for conversion.
```

```
// +kubebuilder:subresource:scale  
:selectorpath=<string>,specpath=<string>  
,statuspath=<string>  
on type  
    enables the "/scale" subresource on a CRD.
```

```
// +kubebuilder:subresource:status on type  
    enables the "/status" subresource on a CRD.
```

CRD Processing

These markers help control how the Kubernetes API server processes API requests involving your custom resources.

See [Generating CRDs](#) for examples.

- ▶ Show Detailed Argument Help

```
//  
+kubebuilder:pruning:PreserveUnknownFields  
on field  
    ▶ PreserveUnknownFields stops the  
    apiserver from pruning fields which are  
    not specified.
```

```
//  
+kubebuilder:validation:XPreserveUnknownFields  
on field  
    ▶ PreserveUnknownFields stops the  
    apiserver from pruning fields which are  
    not specified.
```

```
// +listMapKey:=<string> on field  
    ▶ specifies the keys to map listTypes.
```

```
// +listType:=<string> on field  
    ▶ specifies the type of data-structure  
    that the list represents (map, set,  
    atomic).
```

```
// +mapType:=<string> on field  
    ▶ specifies the level of atomicity of the  
    map; i.e. whether each item in the map is  
    independent of the others, or all fields  
    are treated as a single unit.
```

CRD Validation

These markers modify how the CRD validation schema is produced for the types and fields they modify. Each corresponds roughly to an OpenAPI/JSON schema option.

See [Generating CRDs](#) for examples.

- ▶ Show Detailed Argument Help

```
// +kubebuilder:default:=<any> on field  
    ▶ sets the default value for this field.
```

```
//  
+kubebuilder:validation:EmbeddedResource  
on field  
    ▶ EmbeddedResource marks a fields as  
    an embedded resource with apiVersion,  
    kind and metadata fields.
```

```
// +kubebuilder:validation:Enum:=</>any>  
on field  
    specifies that this (scalar) field is  
    restricted to the *exact* values specified  
    here.
```

```
//  
+kubebuilder:validation:ExclusiveMaximum  
:=bool< on field  
    indicates that the maximum is "up to"  
    but not including that value.
```

```
//  
+kubebuilder:validation:ExclusiveMinimum  
:=bool< on field  
    indicates that the minimum is "up to"  
    but not including that value.
```



Workflow

- ✓ Use sdk to create scaffold project (\$ operator-sdk init)
- ✓ Create scaffold API , CRD, and controller (\$ operator-sdk create api)
- ✓ Update API to have necessary fields for spec and status (\$ edit *_types.go)
- ✓ After above, automatically update generated code for resource - updates interface/implementation for runtime.Object to be a valid Kubernetes Kind (\$ make generate)
 - ✓ updates api/<version>/zz_generated.deepcopy.go
- ✓ Generate CRD manifests (\$ make manifests)
 - ✓ creates config/crd/bases/<item>.yaml
 - ✓ creates config/rbac/role.yaml
- ✓ Update scaffolded controller code to complete controller functionality
 - ✓ Add needed packages and core functions from Kubernetes
 - ✓ Add code to reconcile desired state of CR against observed state
 - ✓ Add code to create any needed K8s objects (Kinds)
 - ❑ Add any additional markers to manage items such as permissions/generate manifests
 - ❑ updates controllers/<item>_controller.go
- ❑ Update CRD manifests (\$ make manifests)
 - ❑ updates config/crd/bases/<item>.yaml
 - ❑ updates config/rbac/role.yaml
- ❑ Execute install target to register CRD with K8s apiserver (\$ make install)
- ❑ Test the operator (\$ make run)
- ❑ Create the docker package (\$ make docker-build)
- ❑ Push image to registry (\$ make docker-push)
- ❑ Deploy operator to cluster (\$ make deploy)

Lab 4 - Specifying Permissions and Adding RBAC Manifests



main.go

- Basic imports
- Scheme - provides mappings between Kinds and Go types
- Basic flags for metrics
- Leader election - only one instance can be in charge (for example if multiple during upgrade)
- Manager instantiation
 - manager keeps tracks of running controllers, sets up shared caches and API client connections
 - runs controllers and webhooks
 - runs until it gets a graceful shutdown signal
 - argument to manager instantiation defines if single or multiple namespaces

```
var (
    scheme = runtime.NewScheme()
    setupLog = ctrl.Log.WithName("setup")
)

func init() {
    utilruntime.Must(clientgoscheme.AddToScheme(scheme))
    utilruntime.Must(roarappv1alpha1.AddToScheme(scheme))
    // +kubebuilder:scaffold:scheme
}

func main() {
    var metricsAddr string
    var enableLeaderElection bool
    flag.StringVar(&metricsAddr, "metrics-addr", ":8080", "The address binds to:")
    flag.BoolVar(&enableLeaderElection, "enable-leader-election", false,
        "Enable leader election for controller manager. "+
            "Enabling this will ensure there is only one active controller")
    flag.Parse()

    ctrl.SetLogger(zap.New(zap.UseDevMode(true)))

    mgr, err := ctrl.NewManager(ctrl.GetConfigOrDie(), ctrl.Options{
        Scheme:                 scheme,
        MetricsBindAddress: metricsAddr,
        Port:                   9443,
        LeaderElection:         enableLeaderElection,
        LeaderElectionID:       "53dff33f.roarapp.com",
    })
    if err != nil {
        setupLog.Error(err, "unable to start manager")
        os.Exit(1)
    }

    if err = (&controllers.RoarAppReconciler{
        Client: mgr.GetClient(),
        Log:   ctrl.Log.WithName("controllers").WithName("RoarApp"),
        Scheme: mgr.GetScheme(),
    }).SetupWithManager(mgr); err != nil {
        setupLog.Error(err, "unable to create controller", "controller", "roarapp")
        os.Exit(1)
    }
    // +kubebuilder:scaffold:builder

    setupLog.Info("starting manager")
    if err := mgr.Start(ctrl.SetupSignalHandler()); err != nil {
        setupLog.Error(err, "problem running manager")
        os.Exit(1)
    }
}
```



What is the manager?

- Manager (package manager component)
 - required to create, start/stop Controllers
 - provides/initializes shared dependencies
 - » schemes, caches, clients,etc.
 - » Scheme - used to resolve runtime objects to group/version/kind
 - executable that wraps 1 or more Controllers
 - » Controllers have to be started by invoking Manager.start
 - can be built & run locally against a remote cluster OR can be run as container in cluster
 - if in container, needs its own Namespace with ServiceAccount and RBAC perms for resources
 - » configs for that automatically created via running make command
 - example:
 - » assume running cluster referenced in ~/.kube/config
 - » build/run locally against cluster defined in ~/.kube/config

```
$ make run
```

then separately create an instance of your resource

```
$ kubectl apply -f config/<file path>
```
- SDK generates code to create Manager
- manager holds Cache and Client to be used with CRUD operations and communicate with API server

```
diyuser3@training1:~/ops$ tree
.
├── bin
│   └── manager
└── config
    ├── certmanager
    │   └── certificate.yaml
    └── kustomization.yaml
```



Operators Scope

- Operators can have two scopes - cluster-scoped & namespaced-scoped
 - cluster-scoped
 - » can watch resources across cluster
 - » requires cluster-scoped permissions (generally a concern)
 - » default for operator-sdk (init)
 - namespaced-scoped
 - » watches resources in a Namespace
 - » best for flexible deployment
 - » permits namespace isolation for monitoring, failure cases, etc.
 - » allows decoupled upgrades



Changing Operator Scope

- Two components for changing scope

- Manager settings

- » By default, Manager does not have any namespace defined
 - » So watches all namespaces
 - » Add namespace or set of namespaces in call to new Manager
 - » Can also use environment variables

```
mgr, err := ctrl.NewManager(ctrl.GetConfigOrDie(), ctrl.Options{
    Scheme:           scheme,
    MetricsBindAddress: metricsAddr,
    Port:              9443,
    LeaderElection:   enableLeaderElection,
    LeaderElectionID: "53dff33f.roarapp.com",
})

mgr, err := ctrl.NewManager(ctrl.GetConfigOrDie(), ctrl.Options{
    Scheme:           scheme,
    MetricsBindAddress: metricsAddr,
    Port:              9443,
    LeaderElection:   enableLeaderElection,
    LeaderElectionID: "53dff33f.roarapp.com",
    Namespace:        "operator-namespace",
})
```

- Restrict roles and permissions

- » Change role.yaml and role_binding.yaml in config/rbac directory
 - » From ClusterRole/ClusterRoleBinding to Role/Rolebinding

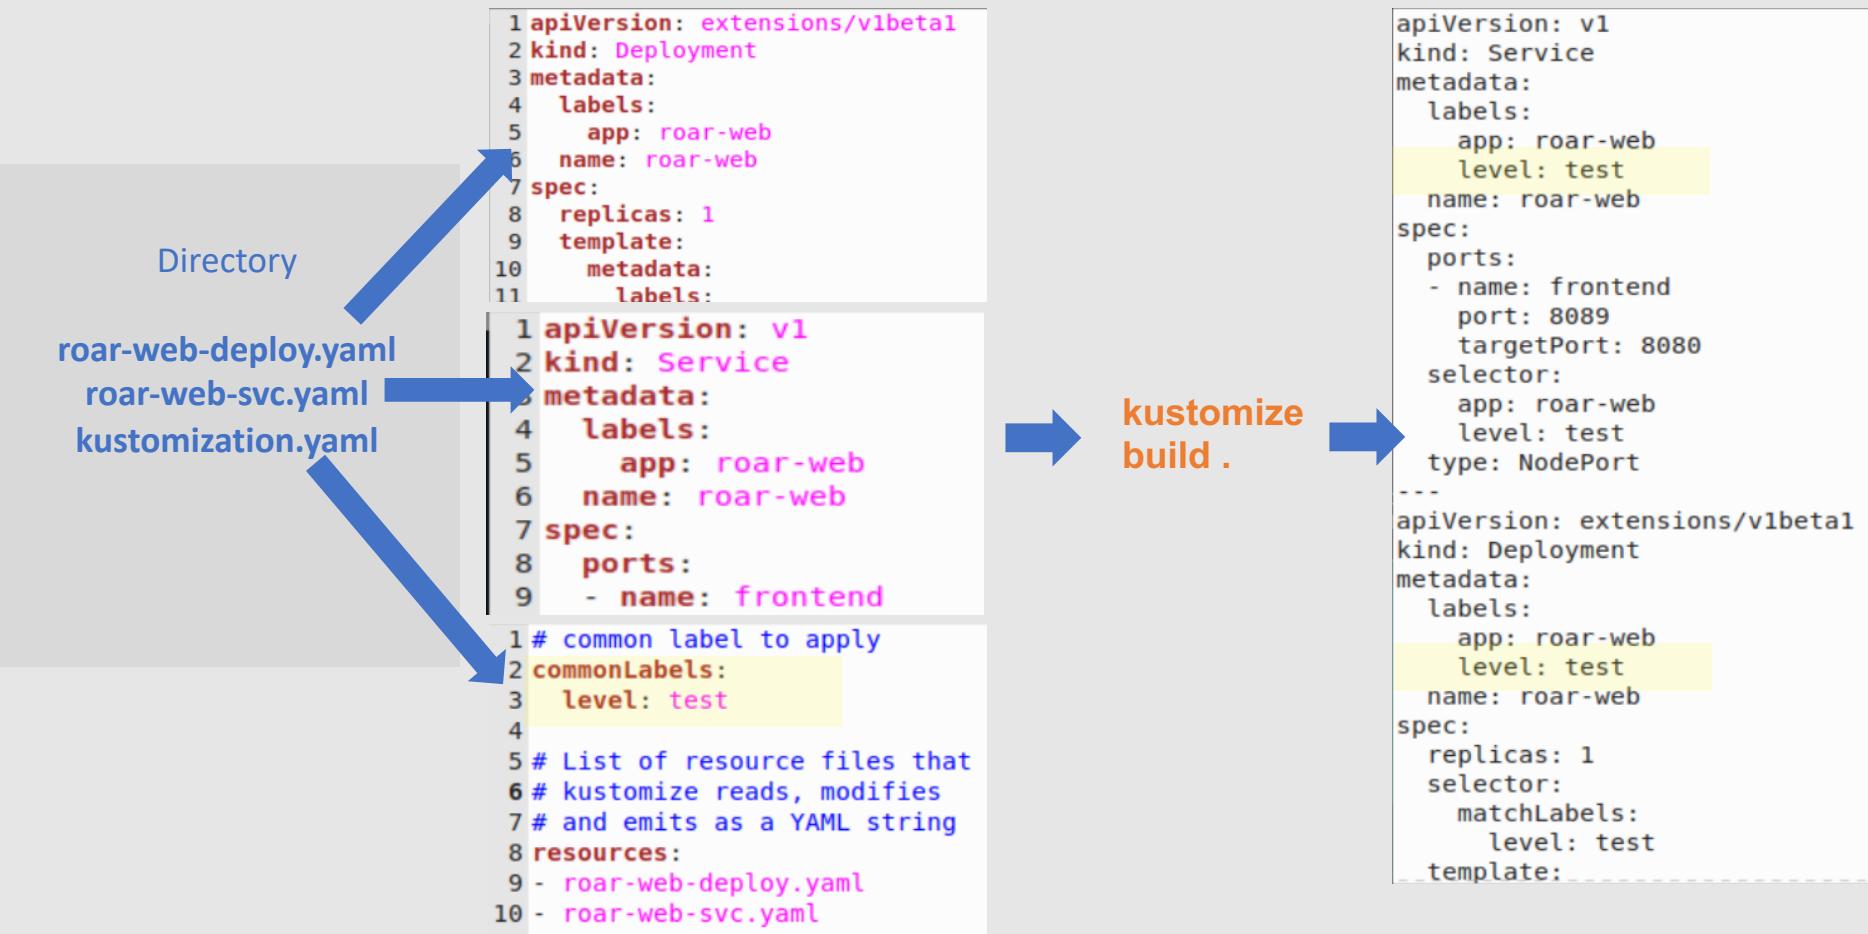


Kustomize

- Kustomize creates modified Kubernetes manifests by “overlaying” declarative specifications on top of existing Kubernetes manifests
- Changes to make/overlay are stored in a separate file called *kustomization.yaml*
- Kustomize reads the *kustomization.yaml* file and the Kubernetes spec files (manifests) and dumps out completed resources to stdout with the changes from *kustomization.yaml* applied “on top of” the standard Kubernetes files
- *kustomize build* command causes files to be processed and output created
- Kustomize leverages the idea of “transformers”
- Transformers are functionality built into Kustomize that “transforms” your Kubernetes manifests given a simple set of declarative rules



Simple Example





Variants

- Common use case is needing multiple variations (variants) of a set of common resources (i.e. dev, stage, prod)
- To do this, Kustomize has concept of “overlay” and “base”
 - Both represented via kustomization.yaml file (includes resources + customizations)
 - Base – declares things variants have in common
 - Overlays – declare the differences
- Variants can also apply “patches”
 - Patch – partial deployment spec – used to update some existing spec as opposed to adding new information



Kustomize Base

```
diyuser3@training1:~/ato-ws/roar-kz/base/db$ kustomize build .
```

```
apiVersion: v1
kind: Service
metadata:
  labels:
    app: roar-db
  1  apiVersion: v1
  2  kind: Service
  3  metadata:
  4  labels:
  5    app: roar-db
  6  name: roar-db
  7 spec:
  8   ports:
  9     - name: mysql
 10    port: 3306
 11    targetPort: 3306
 12 selector:
 13   app: roar-db
```

```
1  apiVersion: v1
2  kind: Service
3  metadata:
4  labels:
5    app: roar-db
6    name: roar-db
7  spec:
8    ports:
9      - name: mysql
10     port: 3306
11     targetPort: 3306
12   selector:
13     app: roar-db
14     path: "/metadata/name"
15     value: "mysql"
16   spec:
17     selector: app: roar-db
```

hat kustomize reads, modifies
ng

Resources to
include

from "roar-db" to "mysql"
ted with the webapp

Patch block = defines
Type of patch
patch in existing text
hierarchy and
replacement value

Item to search for



Running the operator for testing outside of the cluster

```
diyuser3@training1:~/ops$ make run ENABLE_WEBHOOKS=false
go: creating new go.mod: module tmp
go: found sigs.k8s.io/controller-tools/cmd/controller-gen in sigs.k8s.io/controller-tools v0.3.0
/home/diyuser3/go/bin/controller-gen object:headerFile="hack/boilerplate.go.txt" paths=". . ."
go fmt . . .
controllers/roarapp_controller.go
go vet . . .
/home/diyuser3/go/bin/controller-gen "crd:trivialVersions=true" rbac:roleName=manager-role webhook paths=". . ." output:crd:artifacts:config=config/crd/bases
go run ./main.go
2021-02-18T21:12:04.958-0500 INFO controller-runtime.metrics metrics server is starting to listen {"addr": ":8080"}
2021-02-18T21:12:04.958-0500 INFO setup starting manager
2021-02-18T21:12:04.959-0500 INFO controller-runtime.controller Starting EventSource {"controller": "roarapp", "source": "kind source: /, Kind="}
2021-02-18T21:12:04.959-0500 INFO controller-runtime.manager starting metrics server {"path": "/metrics"}
2021-02-18T21:12:05.064-0500 INFO controller-runtime.controller Starting Controller {"controller": "roarapp"}
2021-02-18T21:12:05.064-0500 INFO controller-runtime.controller Starting workers {"controller": "roarapp", "worker count": 1}
```

```
diyuser3@training1:~$ cat ref/roarapp.yaml
apiVersion: roarapp.roarapp.com/v1alpha1
kind: RoarApp
metadata:
  name: roarapp-sample
spec:
  # Add fields here
  replicas: 3
  dbImage: quay.io/bclaster/roar-db:1.0.2
  webImage: quay.io/bclaster/roar-web:1.0.1
```

```
diyuser3@training1:~$ k apply -n op-test -f ref/roarapp.yaml
roarapp.roarapp.roarapp.com/roarapp-sample created
```

```
diyuser3@training1:~$ k get pods -n op-test
```

NAME	READY	STATUS	RESTARTS	AGE
roarapp-sample-pod-32000	2/2	Running	0	47s
roarapp-sample-pod-32001	2/2	Running	0	46s
roarapp-sample-pod-32002	2/2	Running	0	46s

```
2021-02-18T21:12:05.064-0500 INFO controller-runtime.controller Starting workers {"controller": "roarapp", "worker count": 1}
2021-02-18T21:21:18.825-0500 INFO controllers.RoarApp Reconciling instance {"roarapp": "op-test/roarapp-sample"}
2021-02-18T21:21:18.935-0500 INFO controllers.RoarApp Scaling up pods {"roarapp": "op-test/roarapp-sample", "Currently available": 0, "Required replicas": 3}
2021-02-18T21:21:18.992-0500 INFO controllers.RoarApp Reconciling instance {"roarapp": "op-test/roarapp-sample"}
2021-02-18T21:21:19.002-0500 INFO controllers.RoarApp Scaling up pods {"roarapp": "op-test/roarapp-sample", "Currently available": 1, "Required replicas": 3}
2021-02-18T21:21:19.045-0500 INFO controllers.RoarApp Reconciling instance {"roarapp": "op-test/roarapp-sample"}
2021-02-18T21:21:19.053-0500 INFO controllers.RoarApp Scaling up pods {"roarapp": "op-test/roarapp-sample", "Currently available": 2, "Required replicas": 3}
2021-02-18T21:21:19.098-0500 INFO controllers.RoarApp Reconciling instance {"roarapp": "op-test/roarapp-sample"}
2021-02-18T21:21:19.103-0500 DEBUG controller-runtime.controller Successfully Reconciled {"controller": "roarapp", "request": "op-test/roarapp-sample"}
2021-02-18T21:21:19.103-0500 INFO controllers.RoarApp Reconciling instance {"roarapp": "op-test/roarapp-sample"}
2021-02-18T21:21:19.103-0500 DEBUG controller-runtime.controller Successfully Reconciled {"controller": "roarapp", "request": "op-test/roarapp-sample"}
2021-02-18T21:21:19.118-0500 INFO controllers.RoarApp Reconciling instance {"roarapp": "op-test/roarapp-sample"}
2021-02-18T21:21:19.119-0500 DEBUG controller-runtime.controller Successfully Reconciled {"controller": "roarapp", "request": "op-test/roarapp-sample"}
```



Workflow

- ✓ Use sdk to create scaffold project (\$ operator-sdk init)
- ✓ Create scaffold API , CRD, and controller (\$ operator-sdk create api)
- ✓ Update API to have necessary fields for spec and status (\$ edit *_types.go)
- ✓ After above, automatically update generated code for resource - updates interface/implementation for runtime.Object to be a valid Kubernetes Kind (\$ make generate)
 - ✓ updates api/<version>/zz_generated.deepcopy.go
- ✓ Generate CRD manifests (\$ make manifests)
 - ✓ creates config/crd/bases/<item>.yaml
 - ✓ creates config/rbac/role.yaml
- ✓ Update scaffolded controller code to complete controller functionality
 - ✓ Add needed packages and core functions from Kubernetes
 - ✓ Add code to reconcile desired state of CR against observed state
 - ✓ Add code to create any needed K8s objects (Kinds)
 - ✓ Add any additional markers to manage items such as permissions/generate manifests
 - ✓ updates controllers/<item>_controller.go
- ✓ Update CRD manifests (\$ make manifests)
 - ✓ updates config/crd/bases/<item>.yaml
 - ✓ updates config/rbac/role.yaml
- ❑ Execute install target to register CRD with K8s apiserver (\$ make install)
- ❑ Test the operator (\$ make run)
- ❑ Create the docker package (\$ make docker-build)
- ❑ Push image to registry (\$ make docker-push)
- ❑ Deploy operator to cluster (\$ make deploy)

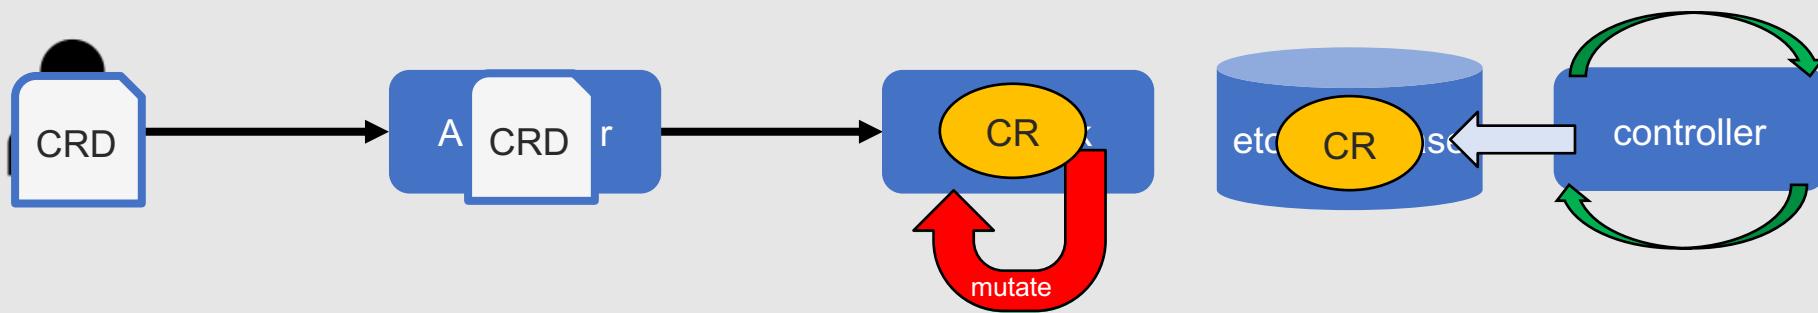


Lab 5 - Building and Testing the Controller



Operators and Webhooks

- Kubernetes api request passes through admission controllers that can invoke webhooks to validate or update api requests
- Webhook is an HTTP "callback" that is registered with the K8s API server
- When a designated event occurs, the K8s server queries registered webhooks and forwards a message
- Can be "mutating" (allowed to modify input objects)
- Can be "validating" (read and validate input objects)



- User creates a CRD and K8s API forwards on to webhook based on registered list.
- Webhook completes any validations/mutations - such as filling in default parameters and/or validating existing ones.
- CR then gets written to database and is available.
- Controller monitors CR in the background and does any needed processing.
- Any updated status resulting from the processing is recorded by controller as status data in the CRD.



Running the Operator in a Cluster

- Create the docker package (`$ make docker-build`)
- Push image to registry (`$ make docker-push`)
- Deploy operator to cluster (`$ make deploy`)

```
diyuser3@training1:~/ops$ k get all -n ops-system
NAME                                         READY   STATUS    RESTARTS   AGE
pod/ops-controller-manager-776bbdbdc7-sls8g   2/2     Running   0          53s

NAME                                TYPE        CLUSTER-IP      EXTERNAL-IP   PORT(S)   AGE
service/ops-controller-manager-metrics-service   ClusterIP   10.104.206.229   <none>       8443/TCP   53s

NAME                           READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/ops-controller-manager   1/1     1            1           53s
replicaset.apps/ops-controller-manager-776bbdbdc7   1         1           1           53s

diyuser3@training1:~/ops$ k logs -n ops-system ops-controller-manager-776bbdbdc7-sls8g -c manager
2021-02-19T03:50:21.806Z     INFO    controller-runtime.metrics      metrics server is starting to listen {"a
ddr": "127.0.0.1:8080"}
2021-02-19T03:50:21.806Z     INFO    setup    starting manager
I0219 03:50:21.806944 [leaderelection.go:242] attempting to acquire leader lease  ops-system/53dff33f.roarap
p.com...
2021-02-19T03:50:21.807Z     INFO    controller-runtime.manager      starting metrics server {"path": "/metrics"
}
I0219 03:50:21.811838 [leaderelection.go:252] successfully acquired lease ops-system/53dff33f.roarapp.com
2021-02-19T03:50:21.811Z     DEBUG   controller-runtime.manager.events      Normal {"object": {"kind": "ConfigM
ap", "namespace": "ops-system", "name": "53dff33f.roarapp.com", "uid": "3f25431e-3d08-4e68-b2e7-11deffe4ca53", "apiVersion
": "v1", "resourceVersion": "19114"}, "reason": "LeaderElection", "message": "ops-controller-manager-776bbdbdc7-sls8g_
41d4206b-fbad-4ad8-bd0d-9875dbdca192 became leader"}
2021-02-19T03:50:21.812Z     INFO    controller-runtime.controller      Starting EventSource {"controller": "roa
rapp", "source": "kind source: /, Kind="}
2021-02-19T03:50:21.913Z     INFO    controller-runtime.controller      Starting Controller {"controller": "roa
rapp"}
2021-02-19T03:50:21.914Z     INFO    controller-runtime.controller      Starting workers {"controller": "roa
```



How Manager container/pod map with main and controller

```
diyuser3@training:~/ops$ k get all -n ops-system
NAME                                     READY   STATUS    RESTARTS   AGE
pod/ops-controller-manager-776bbdbdc7-sls8g   2/2     Running   0          25m

NAME                           TYPE      CLUSTER-IP   EXTERNAL-IP   PORT(S)   AGE
service/ops-controller-manager-metrics-service   ClusterIP  10.104.206.229   <none>    8443/TCP   25m

NAME                           READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/ops-controller-manager   1/1     1           1          25m

NAME                           DESIRED   CURRENT   READY   AGE
replicaset.apps/ops-controller-manager-776bbdbdc7   1         1         1        25m

  "  and so on
  10 RUN go mod download
  11
  12 # Copy the go
  13 COPY main.go
  14 COPY api/
  15 COPY controllers
  16
  17 # Build
  18 RUN CGO_ENABLE=0 go build -o manager main.
  19
  20 # Use distroless
  21 # Refer to https://kubernetes.io/docs/concepts/containers/images/
  22 FROM gcr.io/distroless/base:latest
  23 WORKDIR /
  24 COPY --from=base:rootfs /etc/ssl/certs/ca-certificates.crt /etc/ssl/certs/
  25 USER nonroot:nonroot
  26
  27 ENTRYPOINT ["/manager"]
  28
  29
  30
  31
  32
  33
  34
  35
  36
  37
  38
  39
  40
  41
  42
  43
  44
  45
  46
  47
  48
  49
  50
  51
  52
  53
  54
  55
  56
  57
  58
  59
  60
  61
  62
  63
  64
  65
  66
  67
  68
  69
  70
  71
  72
  73
  74
  75
  76
  77
  78
  79
  80
  81
  82
  83
  84
  85 }

main.go
```



Custom Resource/Types in Production

From API - _types.go

```
// RoarAppSpec defines the desired state of RoarApp
type RoarAppSpec struct {
    // INSERT ADDITIONAL SPEC FIELDS - desired state of cluster
    // Important: Run "make" to regenerate code after modifying this file
    Replicas int32 `json:"replicas"`
    WebImage string `json:"webImage,omitempty"`
    DbImage string `json:"dbImage,omitempty"`
}

// RoarAppStatus defines the observed state of RoarApp
type RoarAppStatus struct {
    // INSERT ADDITIONAL STATUS FIELD - define observed state of cluster
    // Important: Run "make" to regenerate code after modifying this file
    PodNames []string `json:"podNames"`
}
```

```
diyuser3@training1:~$ k get -n ops-system RoarApp
NAME          AGE
roarapp-sample 28m
diyuser3@training1:~$ k get -n ops-system -o yaml RoarApp
apiVersion: roarapp.roarapp.com/v1alpha1
kind: RoarApp
metadata:
  annotations:
    kubectl.kubernetes.io/last-applied-configuration: |
      {"apiVersion":"roarapp.roarapp.com/v1alpha1","kind":"RoarApp","spec":{"replicas":3,"webImage":"quay.io/bclaster/roar-web:1.0.1"}}
    creationTimestamp: "2021-02-19T04:20:31Z"
    generation: 3
  managedFields:
    - apiVersion: roarapp.roarapp.com/v1alpha1
      fieldsType: FieldsV1
      fieldsV1:
        .: {}
        .: {}
      manager: kubectl
      operation: Update
      time: "2021-02-19T04:31:59Z"
    - apiVersion: roarapp.roarapp.com/v1alpha1
      fieldsType: FieldsV1
      fieldsV1:
        .: {}
      manager: manager
      operation: Update
      time: "2021-02-19T04:31:59Z"
    name: roarapp-sample
    namespace: ops-system
    resourceVersion: "21012"
    selfLink: /apis/roarapp.roarapp.com/v1alpha1/namespaces/ops-system/pods/roarapp-sample
    uid: 50cc450-49b7-48c8-a7b3-500c53fab4b0
  spec:
    dbImage: quay.io/bclaster/roar-db:1.0.2
    replicas: 3
    webImage: quay.io/bclaster/roar-web:1.0.1
  status:
    podNames:
      - roarapp-sample-pod-32003
      - roarapp-sample-pod-32004
      - roarapp-sample-pod-32002
```



Workflow

- ✓ Use sdk to create scaffold project (\$ operator-sdk init)
- ✓ Create scaffold API , CRD, and controller (\$ operator-sdk create api)
- ✓ Update API to have necessary fields for spec and status (\$ edit *_types.go)
- ✓ After above, automatically update generated code for resource - updates interface/implementation for runtime.Object to be a valid Kubernetes Kind (\$ make generate)
 - ✓ updates api/<version>/zz_generated.deepcopy.go
- ✓ Generate CRD manifests (\$ make manifests)
 - ✓ creates config/crd/bases/<item>.yaml
 - ✓ creates config/rbac/role.yaml
- ✓ Update scaffolded controller code to complete controller functionality
 - ✓ Add needed packages and core functions from Kubernetes
 - ✓ Add code to reconcile desired state of CR against observed state
 - ✓ Add code to create any needed K8s objects (Kinds)
 - ✓ Add any additional markers to manage items such as permissions/generate manifests
 - ✓ updates controllers/<item>_controller.go
- ✓ Update CRD manifests (\$ make manifests)
 - ✓ updates config/crd/bases/<item>.yaml
 - ✓ updates config/rbac/role.yaml
- ✓ Execute install target to register CRD with K8s apiserver (\$ make install)
- ✓ Test the operator (\$ make run)
- ❑ Create the docker package (\$ make docker-build)
- ❑ Push image to registry (\$ make docker-push)
- ❑ Deploy operator to cluster (\$ make deploy)

Lab 6 - Creating the Operator Image and Deploying it into a Cluster

That's all - thanks!

Professional Git 1st Edition

by Brent Laster (Author)

★★★★★ 3 customer reviews

[Look inside](#) ↓

